

## 8 MySQL (optional – for advanced students)

This lab is optional and can be skipped if you wish – you should only attempt it if you are already experienced with databases and MySQL (or other SQL platforms) or if you are happy to teach it to yourself. If you have already done the Developing Web Applications module you will be able to do this chapter, but if you have not you may want to skip over it for now and come back and attempt it later after you have done the module.

### Learning objectives

The aims of this lab are to:

- Practice using MySQL to create simple databases and tables
- (Optional) Link PHP and MySQL together by creating a dynamic webpage

By the end of this lab session, you should be able to:

- Locate the MySQL daemon, and start, stop and restart it
- Change the behaviour of the MySQL daemon
- Write and execute basic SQL commands to interact with the MySQL server
- Create and view databases and tables using SQL commands
- Insert data into tables using SQL queries

### Introduction

MySQL is the one of the most widely used database management software in the world. It is open source and cross-platform, so you are free to download and install it on your own machine. This lab will introduce you to the basics of running a MySQL server. You are advised to follow this lab carefully and ask questions if you do not understand something. Most importantly, *think about what you are doing!*

Log in to the Slackware machine as root.

### Installing and configuring MySQL

MySQL is preinstalled in Slackware 13.37 by default, so there is no need to compile and install it from source. It does, however, require some more initial configuration compared with Apache before it can be run for the first time.

MySQL comes with some default configuration templates. The active configuration file is located at **/etc/my.cnf**. The template that we will be using in this lab is **my-small.cnf** (there are also **my-medium.cnf**, **my-large.cnf** and **my-huge.cnf**, and you may wish to investigate the differences between them). As root, execute the following command

```
cp /etc/my-small.cnf /etc/my.cnf
```

to copy the template to the active configuration file. Use the following command

```
mysql_install_db
```

to set up the MySQL database for first use. Finally, file and group ownership are fixed by executing:

```
chown -R mysql:mysql /var/lib/mysql
```

## Running MySQL

The following commands will start, stop and restart mysqld, the MySQL daemon, respectively (you may need to chmod `/etc/rc.d/rc.mysqld` first so that it is executable):

```
/etc/rc.d/rc.mysqld start
/etc/rc.d/rc.mysqld stop
/etc/rc.d/rc.mysqld restart
```

If you make any configuration changes to MySQL whilst mysqld is running (including `my.cnf`), you must restart mysqld for those changes to take effect. (Refer to Lab 4 on page 23 if you are stuck.)

You can check if mysqld is running by executing the following command

```
ps -ef | grep mysqld      or      mysqladmin ping
```

Once mysqld is running, we can switch to the MySQL prompt. This is done by executing

```
mysql
```

You should now see the MySQL prompt:

```
mysql >
```

We can now execute SQL commands. All SQL commands must end with a semi-colon (;). To view databases, execute:

```
SHOW DATABASES;
```

and you should see the following output

```
mysql > SHOW DATABASES;
```

```
+-----+
| Database          |
+-----+
| information_schema |
| mysql              |
+-----+
```

We can create a database using the command

```
CREATE DATABASE database_name;
```

and delete a database using the command

```
DROP DATABASE database_name;
```

**WARNING: DON'T DELETE THE mysql DATABASE!** If you do this, you won't be able to log in to MySQL. This can be fixed by executing `mysql_install_db`. **Only delete databases that you have created yourself.**

We can select a database for use using the command

```
USE DATABASE database_name ;
```

Once a database is selected, we can view the tables contained therein using the command

```
SHOW TABLES;
```

**Hint:** MySQL will not parse a query until a semi-colon is typed in. You can therefore break up a long MySQL query into several lines.

---

**Exercise 8.1**

---

1. By executing `SELECT User, Host FROM mysql.user;` find out the number of entries in the **user** table in the **mysql** database.
2. What are the commands to create and delete a table, and to select data from a table? What about inserting data into a table? **Hint:** use the MySQL manual (<http://dev.mysql.com/doc/refman/5.1/en/tutorial.html>)
3. Create a database called **nselab**. Within this database, create a table called **users**, with three fields: **studentid**, **firstname** and **lastname**. Make **studentid** the primary key field. This field cannot be empty for any record and should automatically increase by one each time a new record is added (the first record should have a studentid of 1, the second 2 etc.) The other two fields should be limited to a maximum of 25 characters each. Add two students to this table: Joe Bloggs and Ashley Smith. Devise a plan to test the validation rules (this will involve adding more records) and carry it out.

---

**Exercise 8.2 Optional exercises**

---

1. Create a MySQL user called **adminsec** and grant it full privileges to the **nselab** database only (and no privileges, including viewing, to any other database). You will need to look up the syntax (don't try to do this by editing the **mysql** database directly).
2. Create a web page using PHP that can display the data in the **users** table in the **nselab** database. You should try to show in your logbook that you *understand* what the code does – this could be achieved by annotating the code with comments. (If you are new to PHP, you will need to do some independent research. You may find the following tutorial helpful: <http://www.lynda.com/MySQL-tutorials/PHP-MySQL-Essential-Training/119003-2.html>. You will need to follow the instructions here to log in: <http://web.anglia.ac.uk/it/training/lynda/>.)

## 9 \* Network traffic analysis

**This is an assessed lab. You must write up and submit your answers to this lab for assessment as part of your logbook.**

### Learning objectives

The aims of this lab are to:

- Use shell scripting to carry out network traffic analysis, combining regular expressions and redirection
- Explore “spoofing” a MAC address, DNS lookup and the routing table

By the end of this lab session, you should be able to:

- Understand UNIX commands involving regular expressions and redirection
- Temporarily spoof the MAC address of a Linux system
- Use shell scripting to find the IP address of all machines alive on the same network

### Introduction

Networking is the act of interconnecting machines to form a network so that the machines can interchange information. The most widely used networking stack is TCP/IP, where each node is assigned a unique IP address for identification. There are many parameters in networking, such as subnet mask, route, ports, host names, and so on which require a basic understanding to follow.

Several applications that make use of a network operate by opening and connecting to something called ports, which denote services such as data transfer, remote shell login, and so on. Several interesting management tasks can be performed on a network consisting of many machines. Shell scripts can be used to configure the nodes in a network, test the availability of machines, automate execution of commands at remote hosts, and so on.

Login to the Slackware machine as root.

### Printing the list of network interfaces

Here is a one-line command sequence to print the list of network interfaces available on a system:

```
ifconfig | cut -c-10 | tr -d ' ' | tr -s '\n'
```

```
eth0  
lo  
wlan0
```

The first 10 characters of each line in the `ifconfig` output is reserved for writing names of network interfaces. Hence, we use:

- `cut` to extract the first 10 characters of each line
- `tr -d ' '` to delete every space character in each line
- `tr -s '\n'` to squeeze the `\n` newline character, producing a list of interface names

Depending on how the networking is set up on your Slackware virtual machine, you will normally see either **eth0** or **wlan0**, but not both. **eth0** will be used for the rest of this lab, but if you see **wlan0**, change the commands appropriately.

## Displaying IP addresses

The `ifconfig` command displays details of every active network interface available on the system. However, we can restrict it to a specific interface using:

```
ifconfig interface_name
```

For example:

```
ifconfig eth0

eth0
Link encap:Ethernet
HWaddr 00:1c:bf:87:25:d2      [Hardware (MAC) address]
inet addr:192.168.0.82      [IP address]
Bcast:192.168.3.255        [Broadcast address]
Mask:255.255.252.0         [Subnet mask]
```

In several scripting contexts, we may need to extract any of these addresses from the script for further manipulations. Extracting the IP address is a frequently needed task. In order to extract the IP address from the `ifconfig` output one could use:

```
ifconfig eth0 | egrep -o "inet addr:[^ ]*" | grep -o "[0-9.]*"

192.168.0.82
```

Here, the first command `egrep -o "inet addr:[^ ]*" | grep -o "[0-9.]*"` will print `inet addr:192.168.0.82`. The pattern starts with `inet addr:` and ends with some non-space character sequence (specified by `[^ ]*`). In the next pipe, it prints the character combination of digits and “.”.

## Spoofing the hardware address (MAC address)

In certain circumstances where authentication or filtering of computers on a network are based on the hardware address, we can use hardware address spoofing. The hardware address appears in the `ifconfig` output as `HWaddr 00:1c:bf:87:25:d2`.

We can spoof the hardware address at the software level as follows:

```
ifconfig eth0 hw ether 00:1c:bf:87:25:d5
```

In the preceding command, `00:1c:bf:87:25:d5` is the new MAC address to be assigned. This can be useful when we need to access the Internet through MAC-authenticated service providers that provide access to the Internet for a single machine. However, note that this only lasts until a machine restarts.

## DNS lookup

There are different DNS lookup utilities available from the command line, which will request a DNS server for an IP address resolution. `host` and `nslookup` are two of such DNS lookup utilities. When `host` is executed it will list out all of the IP addresses attached to the domain name. `nslookup` is another command that is similar to `host`, which can be used to query details related to DNS and resolving of names. For example:

```
host google.com

google.com has address 64.252.191.242
[...]
google.com has address 64.252.191.217
```

google.com has IPv6 address 2a00:1450:4009:80d::200e  
 google.com mail is handled by 30 alt2.aspmx.l.google.com.  
 google.com mail is handled by 50 alt4.aspmx.l.google.com.  
 google.com mail is handled by 40 alt3.aspmx.l.google.com.  
 google.com mail is handled by 10 aspmx.l.google.com.  
 google.com mail is handled by 20 alt1.aspmx.l.google.com.

We can also list out all the DNS resource records as follows:

```
nslookup google.com
```

```
Server:    192.168.118.2
Address:   192.168.118.2#53
```

```
Non-authoritative answer:
Name:  google.com
Address: 64.252.191.217
[...]
Name:  google.com
Address: 64.252.191.242
```

The last line in the preceding command-line snippet corresponds to the default name server used for resolution.

Without using the DNS server, it is possible to add a symbolic name to the IP address resolution just by adding entries into the file **/etc/hosts**. In order to add an entry, use the following syntax:

```
echo ip_address symbolic_name >> /etc/hosts
```

For example:

```
echo 192.168.0.9 backupserver >> /etc/hosts
```

After adding this entry, whenever resolution to **backupserver** occurs, it will resolve to 192.168.0.9.

## Showing routing table information

Having more than one network connected with each other is a very common scenario. An example of this is in a college, where different departments may be on separate networks. In this case, when a device on one network wants to communicate with a device on the other network, it needs to go through a device which is common to the two networks. This special device is called a **gateway** and its function is to route packets to and from different networks.

The operating system maintains a table called the **routing table**, which contains the information on how packets are to be forwarded through machines on the network. The routing table can be displayed as follows:

```
route
```

```
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
192.168.0.1	*	255.255.252.0	U	2	0	0	0wlan0
link-local	*	255.255.0.0	U	1000	0	0	0wlan0
default	p4.local	0.0.0.0	UG	0	0	0	0wlan0

Or:

```
route -n
```

#### Kernel IP routing table

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
192.168.0.1	0.0.0.0	255.255.252.0	U	2	0	0wlan0	
169.254.0.0	0.0.0.0	255.255.0.0	U	1000	0	0wlan0	
0.0.0.0	192.168.0.4	0.0.0.0	UG	0	0	0wlan0	

The `-n` option of the `route` command causes all entries to display numerical IP addresses rather than symbolic hostnames.

A default gateway is set as follows:

```
route add default gw ip_address interface_name
```

For example:

```
route add default gw 192.168.0.1 wlan0
```

## Listing all the machines alive on the network

When we deal with a large local area network, we may need to check the availability of other machines in the network. A machine may not be alive due to one of two conditions: either it is not powered on, or due to a problem in the network. By using shell scripting, we can easily find out and report which machines are alive on the network.

---

### Exercise 9.1

---

Create a bash script called **ping.sh** with the following code, then run it. Press [Ctrl] + [z] to quit the program.

```
#!/bin/bash
# Change base address 192.168.0 according to your network.
for ip in 192.168.0.{1..255};
do
    ping $ip -c 2 &> /dev/null;
    if [ $? -eq 0 ];
    then
        echo $ip is alive
    fi
done
```

Think about how the script works before reading the next section.

---

We used the `ping` command to find out the alive machines on the network. We used a `for` loop for iterating through a list of IP addresses generated using the expression `192.168.0.{1..255}`. The `{start..end}` notation will expand and will generate a list of IP addresses, such as `192.168.0.1`, `192.168.0.2`, `192.168.0.3` up to `192.168.0.255`.

`ping $ip -c 2 &> /dev/null` will run a `ping` command to the corresponding IP address in each execution of the loop. The `-c` option is used to restrict the number of echo packets to be sent to a specified number. `&> /dev/null` is used to redirect both `stderr` and `stdout` to `/dev/null` so that it won't be printed on the terminal. Using `$?` we evaluate the exit status. If it is successful, the exit status is 0, else it is non-zero. Hence, the IP addresses which replied to our ping are printed.

In this script, each `ping` command for the IP address is executed one after the other. Even though all the IP addresses are independent of each other, the `ping` command is executed as a sequential program, it takes a delay of sending two echo packets and receiving them or the time-out for a reply for executing the next `ping` command.

# 10 \* Further UNIX tools

**This is an assessed lab. You must write up and submit your answers to this lab for assessment as part of your logbook.**

## Learning objectives

The aims of this lab are to:

- Gain experience in using “classic” UNIX user and system management tools
- Explore symlinks and further file management commands
- Introduce file compression and backup utilities built into UNIX

By the end of this lab session, you should be able to:

- Use tools such as `finger` to find out more information about a UNIX user
- Create symlinks and hard links, and explain the difference between the two
- Compress files into archives, and then extract them

Log in to the Slackware machine as root.

## User and system information

`users` and `who` show the list of users logged into a machine. `who` also shows the terminal they are using and the date they logged in on.

```
users
```

```
bob
```

```
who
```

```
bob      tty1      2014-12-05 19:41
```

`w` gives more detailed information than `who`.

```
w
```

```
USER    TTY    FROM          LOGIN@  IDLE   JCPU   PCPU       WHAT
bob     tty1                08:45   1.00s   0.06s   0.00s      w
```

`finger` also gives detailed information about the list of users logged into a machine, but it can also give detailed information about a specific user (whether or not they are currently logged in).

```
finger ashley
```

```
Login name:      ashley
Registered name: Dr A. Smith      Directory:  /home/ashley
Personal name:    Ashley Smith    Shell:      /bin/bash
Affiliation(s):   Department of Computer Science
Last login Tue Aug 6 18:24 2013 (BST) on /dev/pts/0 from localhost
Mail last read Tue Aug 6 18:36 2013 (BST)
Plan:
To think of a plan.
```



In order to get information about previous boot and user login sessions, use the last command.

last

```
smith      tty2                      Fri Feb  5 10:31    still logged in
bob        tty1                      Fri Feb  5 08:45    still logged in
reboot     system boot 3.19.0-25-generi Fri Feb  5 08:45 - 10:48 (02:02)
bob        tty1                      Thu Feb  4 20:07 - down  (01:48)
reboot     system boot 3.19.0-25-generi Thu Feb  4 20:07 - 21:55 (01:48)

wtmp begins Thu Feb  4 20:07:31 2016
```

last can also be used to obtain information about login sessions for a single user or just the reboot sessions.

last bob

```
bob        tty1                      Fri Feb  5 08:45    still logged in
bob        tty1                      Thu Feb  4 20:07 - down  (01:48)

wtmp begins Thu Feb  4 20:07:31 2016
```

last reboot

```
reboot     system boot 3.19.0-25-generi Fri Feb  5 08:45 - 10:48 (02:02)
reboot     system boot 3.19.0-25-generi Thu Feb  4 20:07 - 21:55 (01:48)

wtmp begins Thu Feb  4 20:07:31 2016
```

Information about failed user logins (if any) can be displayed by executing lastb.

lastb

```
bob        tty1                      Fri Feb  5 10:54 - 10:54 (00:00)
UNKNOWN    tty1                      Fri Feb  5 10:53 - 10:53 (00:00)
root       tty1                      Fri Feb  5 10:52 - 10:52 (00:00)

btmp begins Fri Feb  5 10:57:16 2016
```

In order to see how long the system has been powered on, use the uptime command.

uptime

```
12:33 up 12 days, 20:57, 2 users, load averages: 1.79, 1.84, 1.78
```

---

### Exercise 10.1 User and system information

Before attempting the questions below, you may wish to deliberately reboot the machine and create some failed login attempts so that you have some data to work with.

1. How many login attempts (successful and failed) occurred in the past 48 hours?
  2. How many system reboots occurred in the past 48 hours?
- 

## Symbolic and hard links

In the \* [Email under Linux](#) lab, you were introduced to the mail command. In Slackware 13.37, the program is actually called mailx. Some (older) Linux distributions use nail instead.

So how can we tell whether to use mail, mailx or nail? The command

```
whereis program
```

tells you the path to the binaries in the system, so the command

```
whereis mailx
```

will tell you that mailx is in **/usr/bin/mailx** (and also return a number of other answers). If you do the same for mail, you should find that it returns **/bin/mail**. If you execute

```
ls -l /bin/mail
```

you will see that **/bin/mail** is a symlink (symbolic link) to **/usr/bin/mailx**. Therefore it does not matter if we execute mail or mailx: they are exactly the same. You may wish to try the above for nail as well.

Symlinks are similar to shortcuts in Microsoft Windows and aliases in Mac OS X. Symlinks may be created using the syntax below:

```
ln -s /path/to/original/file /path/to/symlink
```

Hard links may be created as above, but without the -s option.

---

### Exercise 10.2 Symbolic and hard links

---

1. Create a file **~/unixstuff/extra\_file** and a symlink **~/unixstuff/links/extra\_file\_link** which links to **extra\_file** (you may need to create the **links** directory). Use **ls -l** whilst in **~/unixstuff/links/** to check that the symlink has been created.
  2. Edit **extra\_file** and add some text to it. Now open **extra\_file\_link** by executing the following command:  

```
cat ~/unixstuff/links/extra_file_link
```

Do you see the changes you made?
  3. Move **extra\_file** to the **backups** directory (so its location is now **~/unixstuff/backups/extra\_file**).
    - a) What happens to **extra\_file\_link** (if anything)? **Hint:** try opening the symlink using cat, what is the result? Execute **ls -l** whilst in **~/unixstuff/links/**, do you notice anything different?
    - b) Move **extra\_file** back to the **unixstuff** directory – predict what happens to **extra\_file\_link** then test your prediction.
  4. Delete **extra\_file\_link**. What happens to **extra\_file** (if anything – try opening it using cat)?
  5. Recreate the **extra\_file\_link** symlink and delete **extra\_file**. What happens to **extra\_file\_link** (if anything)? See the hint to question 3 (a) if you are stuck.
  6. Delete **extra\_file\_link** and redo questions 1 – 5 above, but this time use hard links instead. Hence explain the differences between symbolic and hard links. You might also wish to do some research to explain why you see these differences.
- 

## File management

Some useful file management commands are:

- **df** (disk free)

The **df** command outputs a report on the disk space available.

- **diff**

**diff** allows you to compare the content of two text files and outputs the differences. The syntax is as follows:

```
diff file1 file2
```

Lines in **file1** are prefixed with < whilst lines in **file2** are prefixed with >.

- **find**

**find** searches through the file system for files matching specified attributes (such as file name, file contents, size). Execute `man find` to see what options are available.

- **touch**

**touch** updates the access and modification date and time of a file to the current time. The syntax is as follows:

```
touch file
```

If **file** does not already exist, it is created with zero contents.

## File compression and backup

UNIX systems usually support a number of utilities for backing up and compressing files. The most useful are:

- **tar** (tape archiver)

**tar** backs up entire directories and files onto a tape device or (more commonly) into a single disk file known as an archive. An archive is a file that contains other files plus information about them, such as their filename, owner, timestamps, and access permissions. **tar** does not perform any compression by default.

To create a disk file **tar** archive, use

```
tar -cvf archivename filename
```

where *archivename* will usually have a **.tar** extension. Here the **-c** option means create, **-v** means verbose (output filenames as they are archived), and **-f** means file. To list the contents of a **tar** archive, use

```
tar -tvf archivename
```

To restore files from a **tar** archive, use

```
tar -xvf archivename
```

- **cpio**

**cpio** is another facility for creating and reading archives. Unlike **tar**, **cpio** doesn't automatically archive the contents of directories, so it's common to combine **cpio** with **find** when creating an archive:

```
find . -print -depth | cpio -ov -H tar > archivename
```

This will take all the files in the current directory and the directories below and place them in an archive called *archivename* . The **-depth** option controls the order in which the filenames are produced and is recommended to prevent problems with directory permissions when doing a restore. The **-o** option creates the archive, the **-v** option prints the names of the files archived as they are added and the **-H** option specifies

an archive format type (in this case it creates a tar archive). Another common archive type is `crc`, a portable format with a checksum for error control.

To list the contents of a `cpio` archive, use

```
cpio -tv < archivename
```

To restore files, use:

```
cpio -idv < archivename
```

Here the `-d` option will create directories as necessary. To force `cpio` to extract files on top of files of the same name that already exist (and have the same or later modification time), use the `-u` option.

- **compress and gzip**

`compress` and `gzip` are utilities for compressing and decompressing individual files (which may be or may not be archive files). To compress files, use:

```
compress filename      or      gzip filename
```

In each case, *filename* will be deleted and replaced by a compressed file called **filename.Z** or **filename.gz**. To reverse the compression process, use:

```
compress -d filename    or      gzip -d filename
```

`zcat` can be used to read gzipped text files without uncompressing them first. The output can be piped to `less` if the text file is very long.

---

**Exercise 10.3 File compression and backup**

---

1. Archive the contents of your home directory (including any subdirectories) using `tar` and `cpio`.
  2. Compress the tar archive with `compress`, and the `cpio` archive with `gzip`.
  3. Now extract their contents.
-